

Part-A

1. OS acts as an intermediary between hardware and software.

- 2) - It manages resources (CPU, memory, I/O) efficiently.
- Provides essential services like process scheduling, file systems, and security.
- Provides essential services like process scheduling, file systems, and security.

3) - Real-Time Operating System (RTOS) is best.

- RTOS ensures deterministic and timely response, which is crucial for health monitoring.
- light weight, power-efficient, and can handle sensor inputs reliably.

4) - Avoids Monolithic kernel.

- Because it has all services in one large module, leading to:
 - Large codebase and poor modularity.
 - Higher chance of bugs affecting the whole system.
 - Difficult debugging and context switching.

5) - Refute.

- OS structure affects stability, performance, maintainability, and security.
- Poor structure can cause bottlenecks, crashes, or vulnerability exploits even if processes run.
- Structured design ensures modularity and scalability.

6) 1) PCB stores process state (registers, program counter, flags).

- Analyzing PCB shows if registers or states are misinitialized during switch.

ii) Context switching involves:

- Saving the current process state in its PCB.
- Loading the next process's state from its PCB.
- Updating CPU registers, program counter, and memory mappings.

iii) Use a blocking synchronous system call.

- Blocking ensures the process waits for I/O allocation to complete.
- Synchronous ensures tasks proceed only after I/O allocation succeeds (avoids race conditions).

Part-B

- Q:
- Save State time = 2ms
 - Load State time = 3ms
 - Scheduler Overhead = 1ms

a) Total context switching time (T_{-cs}):

$$T_{-cs} = (\text{Save State}) + (\text{Load State}) + (\text{Scheduler overhead}) \\ = 2\text{ms} + 3\text{ms} + 1\text{ms}$$

$$T_{-cs} = 6\text{ms}$$

b) - This overhead adds to every switch, reducing effective CPU time for actual processes.

- Frequent context switches \rightarrow lower multitasking efficiency

Q1. Given - Total execution time = 40 sec

If the task is perfectly parallelizable and there are N threads \rightarrow

$$T_{\text{parallel}} = T_{\text{single}}/N$$

• For $N = 2$ threads $\rightarrow T = 40/2 = 20 \text{ sec}$

• For $N = 4$ threads $\rightarrow T = 40/4 = 10 \text{ sec}$

• For $N = 8$ threads $\rightarrow T = 40/8 = 5 \text{ secs}$

Q2. Realistic : Amdahl's law

If a function ϕ of the program is parallelizable and $(1-p)$ is serial, max speedup (SCN) using N processors/threads is:

$$SCN = \frac{1}{(1-p) + SP/N}$$

$$T_{\text{parallel}} = \frac{T_{\text{single}}}{SCN}$$

80% parallelizable $\rightarrow p = 0.8$, Serial = 0.2

Compute speedup for $N=4$

$$SCN = \frac{1}{(0.2+0.8)} = \frac{1}{0.2+0.8} = \frac{1}{0.4} = 0.25$$

$T_{\text{parallel}} = 40/2.5 = 16 \text{ sec}$

Compare

• Ideal 4-thread time = 10 sec

• Realistic ($p = 0.8$) 4-thread time = 16 sec

Explanation for benefits

- Thread reduces total time when substantial parts are parallelizable.
- Overheads (synchronization, context switching, Cache coherence) reduce ideal gains - use concurrency where it truly benefits.

<u>Q:</u> Process	Burst Time
P ₁	5
P ₂	3
P ₃	8
P ₄	6

All arrive at time 0.

a) ~~FCFS (first come, first served)~~

Order: P₁ → P₂ → P₃ → P₄

- P₁: start 0, burst 5 → completes at $0 + 5 = 5$
- P₂: starts at 5, burst 3 → completes at $5 + 3 = 8$
- P₃: starts at 8, burst 8 → completes at $8 + 8 = 16$
- P₄: starts at 16, burst 6 → completes at $16 + 6 = 22$.

Grant timeline showing [start - end] for each:

|P₁ [0-5]| P₂ [5-8] | P₃ [8-16] | P₄ [16-22]|

2. Non-preemptive SJF (shortest Job first)

Order by burst ascending: P₂(3), P₁(5), P₄(6), P₃

- P₂: starts at 0 → completes at $0 + 3 = 3$

- P₁: starts at 3 → completes at $3 + 5 = 8$

• P₄: start 8 → completes at $8+6 = 14$

• P₃: start 14 → completes $14+8 = 22$.

Grantt:

$|P_2[0-3]|P_1[3-8]|P_4[8-14]|P_3[14-22]|$

2'. Round Robin (Quantum Q = 4 ms)

~~Simulate~~ Simulate cycles (all arrive at 0). Keep remaining times:

Initial remaining = ~~(P₁: 5, P₂: 3, P₃: 8, P₄: 6)~~

• Cycle 1:

- Time 0-4: P₁ runs for 4 (remaining P₁ = 5-4 = 1). Time now = 4.

- Time 4-7: P₂ needs 3 (≤ 4), runs 3, completes at 7. (P₂ remaining = 0).

- Time 7-11: P₃ runs 4 (remaining = 8-4 = 4), Time = 11.

- Time 11-15: P₄ runs 4 (remaining = 6-4 = 2). Time = 15.

Remaining after cycle 1: P₁ : 1, P₂ : 0 (done), P₃: 4, P₄: 2.

• Cycle 2

• Time 15-16: P₁ runs 1, completes at 16.

• Time 16-20: P₃ runs 4, completes at 20.

• Time 20-22: P₄ runs 2, completes at 22.

Grantt (with time slices)

$|P_1[0-4]|P_2[4-7]|P_3[7-11]|P_4[11-15]|P_3[15-16]|$

$P_3[16-20]|P_4[20-22]|$

- b)
- Turnaround Time (TAT) = completion time - arrival time. (arrival = of arrival)
 - Waiting Time (WT) = Turnaround Time - Burst time.

FIFO

$$\bullet \text{WT} : (0+5+8+16) = 29 \rightarrow 29/4 = 7.25$$

$$\bullet \text{TAT} : (5+8+16+22) = 51 \rightarrow 51/4 = 12.75$$

SJF

~~$$\bullet \text{WT} : (3+0+14+8) = 25 \rightarrow 25/4 = 6.25$$~~

~~$$\bullet \text{TAT} : (8+3+22+14) = 47 \rightarrow 47/4 = 11.75$$~~

RR

~~$$\bullet \text{Completion} : P_1 = 16, P_2 = 7, P_3 = 20, P_4 = 22$$~~

~~$$\bullet \text{WT} = \text{TAT} - \text{Burst}$$~~

~~$$\bullet \text{TAT} : (16+7+20+22) = 65 \Rightarrow 65/4 = 16.25$$~~

~~$$\bullet \text{WT} : (11+4+12+16) = 43 \rightarrow 43/4 = 10.75$$~~

c) - SJF gives lowest average waiting time and turnaround time.

- Hence it balances throughput and turnaround time among the three.

9) Cloud migration OS architecture

Architecture:

• Microkernel - because:

• Minimal core increases efficiency.

• Services run in user mode, improving fault isolation.

- Easy scalability by adding services as modules.

Virtual Machines Role:

- Provide isolated environments for each service.
- Enable resource pooling and dynamic allocation.
- Simplify system management and migration without affecting others.

(ii) Smart Home System Scheduling

Process scheduling + IPC:

- OS assigns higher priority to critical tasks (intrusion alerts).
- lower priority for less critical tasks (light, thermostat).
- IPC (signals, message queues) ensure quick communication between processes.

Suitable algo:

- Priority Scheduling (Preemptive) → ensures ~~of~~ ~~critical tasks~~ ~~in~~ ~~dead~~
- Earliest Deadline First (EDF) → good for real-time constraints.
- Round Robin (for non-critical tasks) to ensure fairness.