

OS Assignment - 2

Ayush Gupta
2301010206
BTech CSE - D

Part-A

- 1) Address translation with multiple processes.
- Each process has its own page table.
 - MMU + TLB translate virtual to physical addresses.
 - On context switch, OS loads new page table base; ASID to avoid TLB flush.
 - Page faults handled by OS to bring pages from disk.
- 2) Layout causing both internal & external fragmentation + solutions.
- External: variable-size allocations create small scattered holes.
 - Internal: paging wastes space in last page (e.g. 4000B \rightarrow 4+1 pages).
 - Fixes: paging (removes external), segmentation + paging hybrid, buddy & slab allocators, coalescing free lists, memory compaction (ram), variable page sizes, GC compaction.
- 3) Paging-based model & trade-offs
- Fixed-size pages (e.g. 4KB), multi-level page tables, TLB cache, swapping.
 - Small pages: less internal waste, more page table entries, more TLB misses.
 - Large pages: fewer TLB misses, smaller page tables, more internal waste.

- Multi-level tables reduce memory use but slow on TLB miss.

4) OS-Hardware interaction in virtual memory

- Hardware: MMU, TLB, page table base register (CR3), ASIDs, page-fault trap.
- Steps: TLB hit \rightarrow direct; TLB miss \rightarrow page-table walk; if not present \rightarrow page fault - OS loads page \rightarrow update PTE \rightarrow resume.

5. Given: Virtual address width = 16 bits \rightarrow virtual address space size = 2^{16} bytes = 65536 bytes.

Page size = 1KB = 1024 bytes.

size per page - table entry (PTE) = 2 bytes.

a) No. of virtual pages:

$$\text{No. of pages} = \frac{\text{virtual address space}}{\text{page size}} = \frac{65536}{1024}$$

= 64 virtual pages.

b) Page table size (bytes) = No. of pages \times PTE size
 $= 64 \times 2$
 $= 128 \text{ bytes.}$

Part-B

6. ~~Memory~~ Memory allocation simulation

Given: free memory = 1000 KB. Process - $P_1 = 212 \text{ KB}$,
 $P_2 = 417 \text{ KB}$, $P_3 = 112 \text{ KB}$, $P_4 = 426 \text{ KB}$.

First - Fit

1. Start hole = 1000 KB
2. Place P_1 (212 KB) \rightarrow remaining hole = $1000 - 212 = 788 \text{ KB}$
3. Place P_2 (417 KB) \rightarrow remaining hole = $788 - 417 = 371 \text{ KB}$
4. Place P_3 (112 KB) \rightarrow remaining hole = $371 - 112 = 259 \text{ KB}$
5. Try P_4 (426 KB) $\rightarrow 259 \text{ KB} < 426 \text{ KB} \rightarrow$ can't place.

Result (First - Fit): allocated P_1, P_2, P_3 ; free leftover.
= 259 KB; P_4 unallocated

Best-fit: same sequence here (only one hole at each step) \rightarrow 259 KB leftover (P_4 unallocated).

Worst-fit: same result \rightarrow 259 KB leftover.

Conclusion: All three produces identical results for this sequence. ~~259 KB~~ 259 KB unused, P_4 not placed.

7. Page Replacement

Reference string: - 7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 2
frames: 3

I'll give concise step outcomes (frame contents after each reference shows when a change happens).

FIFO (replace oldest)

Step through references - faults occur when page not in frames.

- After references and replacements in FIFO page faults = 10

Optimal (replace page whose next use is farthest)

Using future knowledge, optimal choose best victim.

- optimal page faults = 7

LRU (replace least - recently used)

LRU page faults = 9

Algorithm	Page faults
FIFO	10
Optimal	7
LRU	9

8. Dirty page write overhead.

Given: disk WT = 10ms/page, memory write = 10ms/page, dirty fraction = 30%, pages replaced = 1000

Step-1 - No. of dirty pages:

$$1000 \times 0.3 = 300 \text{ pages.}$$

Step-2 - Disk wait time total:

$$300 \times 10 = 3000 \text{ ms} = 3 \text{ secs.}$$

Step-3 Memory write time (for comparison):

$$300 \times 100 \text{ ns} = 30000 \text{ ns} = 0.00003 \text{ s (negligible)}$$

Extra overhead $\sim 0.3 \text{ sec}$. due to dirty ~~page~~ page
disk writes for 1000 replacements (30% dirty)

9. Autonomous vehicle memory management

- Working set model: give real-time tasks (object detection) enough frames - their working set \rightarrow prevent thrashing.
- Policies:
 - lock critical pages in RAM (mlock), use wlock on PFT control.
 - Reclaim from low-priority tasks if critical has fault rate rises.
 - Use memory compression + quotas (cgroups) for infotainment.
 - Prioritized async writes to avoid blocking real-time tasks.